

Enterprise AI-Powered Customer Analytics & Strategic Insight Framework

Churn Prediction, Revenue Forecasting, Behavioural Segmentation and LLM-Generated Executive Intelligence
for a SaaS / CRM Platform (Salesforce / Zoho style)

Attribute	Detail
Domain	SaaS / CRM Analytics / Customer Intelligence
Dataset	Star schema: 4 fact tables, 3 dimension tables
Modelling grain	5,000 customers (one row each)
Features engineered	180 modelling features
Classification target	Churn (binary)
Regression target	Next_Quarter_Revenue_USD (continuous)
Report date	29 July 2026

1. Executive Summary

This project delivers an end-to-end customer intelligence system for a subscription SaaS CRM platform. Four relational fact tables and three dimension tables were consolidated into a single customer-level table of 5,000 accounts and 180 engineered features. Three machine-learning models were trained on that table, their outputs fused into a standardised per-account intelligence profile, and a large language model used to translate those profiles into executive briefings.

The churn classifier achieves a hold-out ROC-AUC of **0.9023** and an F1 of **0.6020** on the minority churn class, after withholding four record-count features identified as target leakage. The revenue model explains **96.7%** of next-quarter revenue variance with an RMSE of **\$1,363**. Segmentation resolves the portfolio into **2** behavioural clusters whose churn rates differ by a factor of roughly seven. Fusing these outputs identifies **\$697,827** of next-quarter revenue sitting in accounts with elevated churn probability.

2. Business Problem to ML Task Mapping

Every technical component is anchored to a measurable business objective, so model performance translates directly into commercial outcomes.

Business Problem	ML Task	Output	Business KPI Served
Customers cancelling subscriptions	Binary classification (XGBoost)	Churn probability 0-1	Retention rate, customer lifetime value
Unreliable revenue planning	Regression (GradientBoosting)	Next-quarter revenue (USD)	Forecast accuracy, quota setting
Undifferentiated marketing	Clustering (KMeans + PCA)	Behavioural cluster ID	Campaign conversion, personalisation
Unclear account prioritisation	Feature fusion / scoring	Risk, potential, priority scores	Revenue at risk, CSM coverage
Technical output not board-ready	LLM narrative generation	Executive briefing text	Decision latency, reporting effort

3. Methodology

3.1 Data engineering

- Fact tables aggregated from their native grain to customer grain: transactions (54,858 rows), monthly usage snapshots (131,268 rows) and engagement events (54,822 rows).
- Dimension attributes joined on Geo_ID, Industry_ID and Product_ID. Duplicate dimension keys were removed first — dim_geography ships five duplicate Geo_ID rows that would otherwise fan out the customer grain.
- 150 duplicate customer records removed, leaving 5,000 unique accounts.
- Missing values imputed with the median (numerical) and mode (categorical); zero missing cells remain.

3.2 Feature engineering

- **Tenure:** contract duration, tenure in days and years, days to contract end.
- **Usage intensity:** sessions per active user weighted by session length, DAU/MAU ratio, API calls per licence, error rate per session.
- **Engagement ratios:** events per tenure month, escalated ticket ratio, support dependency, sentiment-based engagement quality score.
- **Revenue:** quarterly revenue mean, volatility, latest value and growth ratio, average monthly revenue, CLV approximation, payment reliability.

3.3 Encoding and class balance

Binary and high-cardinality categoricals were label encoded; low-cardinality nominals were one-hot encoded, producing 190 columns. The churn target is imbalanced at roughly 11.3% positives. SMOTE was applied **only to the training split after the stratified train/test split**, so no synthetic record can reach the hold-out set.

4. Model Performance

4.1 Churn classification

Metric	Value	Business Interpretation
ROC-AUC (hold-out)	0.9023	Separates churners from retained accounts reliably
F1-score (churn class)	0.6020	Balanced precision and recall on the minority class
Precision (churn class)	0.7108	Share of flagged accounts that genuinely churn
Recall (churn class)	0.5221	Share of actual churners the model catches
Cross-validated ROC-AUC	0.9032 $\pm$ 0.0152	SMOTE refitted inside each fold - no resampling leak
Baseline GradientBoosting ROC-AUC	0.8689	XGBoost selected as the stronger model
Features used	180	After withholding leakage-prone columns

The cross-validated figure is produced by an imbalanced-learn pipeline in which SMOTE is refitted on the training portion of every fold. Cross-validating a model already fitted on pre-balanced data would be optimistic, because synthetic rows derived from a validation record can reach the training folds; that earlier approach reported 0.9954 against the honest 0.9032.

4.1.1 Target-leakage ablation

Four record-count features (transaction count, renewal count, active quarters, months of usage observed) fall simply because a churned customer stops generating rows, so their value encodes the outcome rather than predicting it. They are withheld from the shipped model. Engagement counts are deliberately retained: they correlate *positively* with churn (distressed customers raise more tickets), which is genuine leading signal. The table below states the measured cost of that decision.

Variant	Features	ROC-AUC	F1	Recall
Shipped model (leak-free)	180	0.9023	0.6020	0.5221
Contaminated variant (ablation only)	184	0.9455	0.6881	0.6637

Excluding the artefacts costs 0.0432 ROC-AUC and 0.1416 recall. The lower figure is the honest one: the contaminated variant scores well partly by reading the outcome it is meant to forecast.

4.2 Explainable AI — SHAP churn drivers

Rank	Feature	Mean SHAP	Business Meaning
1	Renewal_Risk_Flag_Low	1.3744	CRM renewal-risk flag set to Low
2	txn_total_revenue_usd	0.7044	Model-derived behavioural signal
3	txn_total_billed_usd	0.6566	Model-derived behavioural signal
4	txn_mean_revenue_usd	0.5815	Model-derived behavioural signal
5	escalated_ticket_ratio	0.3597	Share of support tickets escalated
6	txn_mean_payment_delay_days	0.3192	Average invoice payment delay
7	Relative_Churn_Risk_Medium	0.3099	Industry-level baseline churn risk
8	Relative_Churn_Risk_Low	0.2769	Industry-level baseline churn risk
9	txn_failed_payment_rate	0.2650	Model-derived behavioural signal
10	usage_mean_license_utilization_pct	0.2586	Licences actually used vs purchased

4.3 Revenue forecasting

Metric	Value	Business Interpretation
Selected model	gradient_boosting	Chosen over the alternative on hold-out RMSE
RMSE	\$1,363.28	Typical forecast error, penalising large misses
MAE	\$692.37	Average absolute error per account per quarter
R ²	0.9674	Share of revenue variance explained
Adjusted R ²	0.9600	Corrected for the number of features used
Cross-validated R ²	0.9557 ± 0.0105	Stability across five folds

4.4 Behavioural segmentation

Metric	Value	Assessment
Clusters selected (k)	2	Chosen by best silhouette across k = 2 to 10
Silhouette score	0.3760	Below the 0.60 project target
Davies-Bouldin index	1.1886	Lower is better; indicates moderate compactness
Calinski-Harabasz	2,723.4	Higher is better; confirms real separation
PCA components	6	90% of behavioural variance retained

Honest assessment: the 0.60 silhouette target was not met. The best achievable score was 0.3760 at k=2. SaaS behavioural features are continuous and overlapping, so accounts form a gradient rather than well-separated spheres, and no k between 2 and 10 approached the target. The clusters remain commercially useful: churn rates differ roughly sevenfold between them, which is the property the business actually needs.

5. Customer Segment Profiles

Cluster	Accounts	Mean MAU	Adoption %	Avg Monthly Rev	Churn Rate
0	1,401	46	30.2%	\$576	30.0%
1	3,599	202	73.3%	\$2,127	4.0%

Cluster 0 (*Disengaged At-Risk Accounts*) combines low usage, low adoption and heavy support dependency with a materially higher churn rate. Cluster 1 (*Loyal High-Value Accounts*) shows roughly four times the monthly active users, more than double the feature adoption, and near-negligible churn. Retention spend belongs in cluster 0; expansion and upsell motions belong in cluster 1.

6. Unified Customer Intelligence Layer

The fusion layer scores every account with all three models and emits a standardised JSON profile containing churn probability and risk band, next-quarter revenue forecast, revenue at risk, cluster assignment and label, percentile-normalised risk / potential / priority scores, and the top five per-account SHAP drivers with direction of effect.

Portfolio Indicator	Value
Accounts profiled	5,000
High-risk accounts	516
Medium-risk accounts	61
Low-risk accounts	4,423
Total forecast revenue at risk	\$697,827
Mean churn probability	11.5%
Top 5 accounts, revenue at risk	\$53,087

7. LLM-Based Executive Insight Generation

Structured intelligence profiles are passed to a Gemini model under a system instruction that forbids invented metrics and requires every claim to trace back to a supplied value. Each briefing covers account risk level, revenue impact, behavioural segment interpretation, prioritised strategic recommendations with owner and timeframe, and a targeted renewal plus upsell play. Five briefings were generated for the accounts carrying the highest forecast revenue at risk and exported to reports/executive_insights.md.

Implementation note: the originally specified gemini-pro model has been retired from the API. The pipeline probes current models at runtime and uses the first one the credential can serve.

8. Limitations and Future Work

- **Residual leakage risk.** The four clearest temporal artefacts are now withheld and the cost measured (section 4.1.1), but other aggregates such as lifetime revenue also scale with how long an account survived. A production rebuild should compute every feature from a fixed observation window that closes before the prediction date.
- **Recall ceiling.** The leak-free model catches 52.2% of churners at 71.1% precision. Raising recall means lowering the decision threshold and accepting more false positives; the right operating point depends on the cost of a retention outreach versus a lost account.
- **Segmentation separation.** Silhouette of 0.376 falls short of the 0.60 target; density-based or model-based clustering may suit gradient-shaped behavioural data better.

- **Static snapshot.** The system scores a point-in-time extract. Production deployment needs scheduled re-scoring, drift monitoring and retraining triggers.
- **Insight validation.** LLM briefings are constrained by prompt design but not yet automatically checked; a factuality validator comparing generated numbers against the source profile is the natural next step.

9. Deliverables

Artefact	Location
Preprocessing pipeline	src/preprocessing.py
Model training	src/train_models.py
Intelligence fusion layer	src/fusion_layer.py
LLM insight generation	src/llm_insights.py
PDF report generator	src/generate_report.py
Executed analysis notebook	notebooks/EDA_and_Modeling.ipynb
Serialised models	models/*.pkl
Processed dataset	data/processed/processed_customer_data.csv
Customer intelligence profiles	data/processed/customer_profiles.json
Executive briefings	reports/executive_insights.md
Model metrics	reports/model_metrics.json
This report	reports/Capstone_Final_Report.pdf